

# PROJECT ON INSURANCE:

## PROBLEM STATEMENT:

WHICH MODEL SUITS FOR THE INSURANCE PROBLEM

```
In [119]: import numpy as np
import pandas as pd
import seaborn as sns
from scipy import stats
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn import metrics
from sklearn.linear_model import Lasso, LassoCV
from sklearn.linear_model import Ridge, RidgeCV
from sklearn.preprocessing import StandardScaler
```

## DATA COLLECTION

## READ THE DATA

```
In [120]: df=pd.read_csv(r"C:\Users\HP\Downloads\insurance.csv")
df
```

Out[120]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

## DATA CLEANING

```
In [121]: df.columns
```

Out[121]: Index(['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges'], dtype='object')

```
In [122]: df.describe()
```

```
Out[122]:
```

	age	bmi	children	charges
count	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	30.663397	1.094918	13270.422265
std	14.049960	6.098187	1.205493	12110.011237
min	18.000000	15.960000	0.000000	1121.873900
25%	27.000000	26.296250	0.000000	4740.287150
50%	39.000000	30.400000	1.000000	9382.033000
75%	51.000000	34.693750	2.000000	16639.912515
max	64.000000	53.130000	5.000000	63770.428010

```
In [123]: df.tail()
```

```
Out[123]:
```

	age	sex	bmi	children	smoker	region	charges
1333	50	male	30.97	3	no	northwest	10600.5483
1334	18	female	31.92	0	no	northeast	2205.9808
1335	18	female	36.85	0	no	southeast	1629.8335
1336	21	female	25.80	0	no	southwest	2007.9450
1337	61	female	29.07	0	yes	northwest	29141.3603

```
In [124]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

```
In [125]: df.shape
```

```
Out[125]: (1338, 7)
```

## EXPLORATORY DATA ANALYSIS

### DATA PREPROCESSING

```
In [126]: df.isnull().any()
```

```
Out[126]: age         False
sex           False
bmi           False
children      False
smoker        False
region        False
charges       False
dtype: bool
```

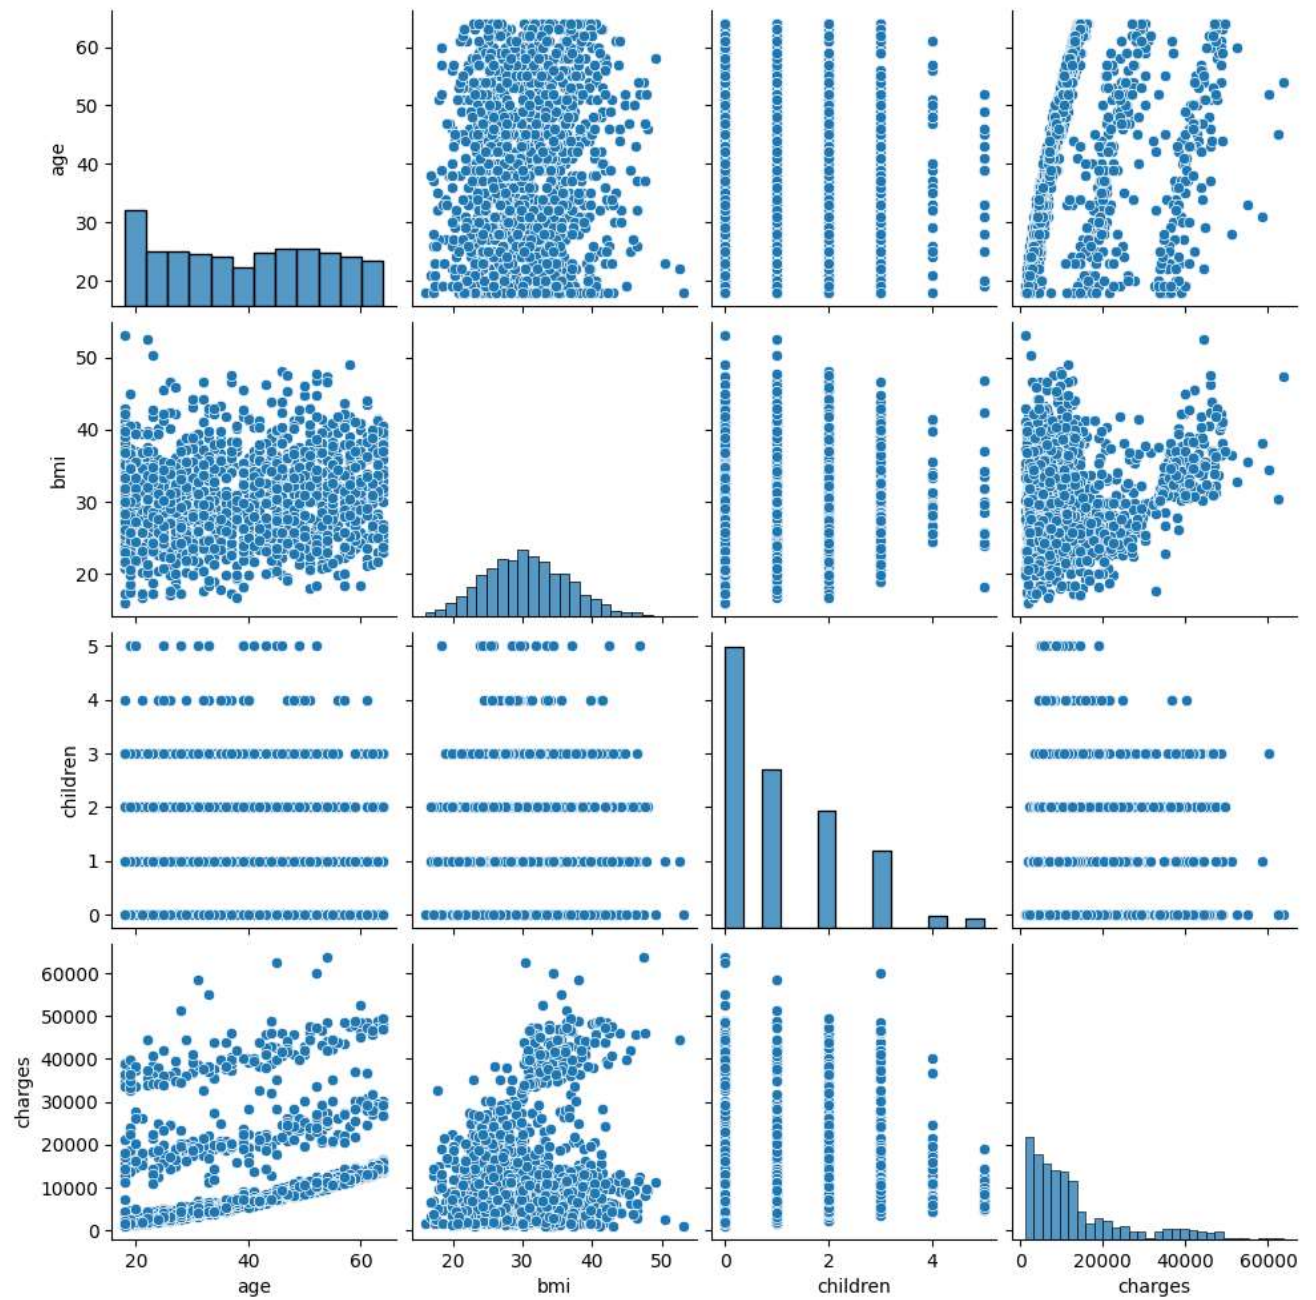
HERE THERE IS NO NULL VALUE.SO THIS IS NOT NECESSARY TO ADJUST THE VALUES

### DATA VISUALIZATON

```
In [127]: #plots
```


```
In [128]: sns.pairplot(df)
```

```
Out[128]: <seaborn.axisgrid.PairGrid at 0x21f20427040>
```



```
In [129]: Insurancedf=df[['age','bmi', 'children','charges']]
sns.heatmap(Insurancedf.corr())
```

Out[129]: <Axes: >



## Feature Scaling:splitting data into train and test data sets

```
In [130]: x=np.array(df['age']).reshape(-1,1)
y=np.array(df['charges']).reshape(-1,1)
```

```
In [131]: from sklearn.linear_model import Ridge,RidgeCV,Lasso
from sklearn.preprocessing import StandardScaler
```

## DATA MODELING:

first we go through REGRESSION models

### 1. LINEAR REGRESSION

```
In [113]: #1.TO MAKE CONVERSIONS TO HAVE SAME DATA TYPE TO FIND REGRESSION MODEL
```

```
In [134]: A={"smoker":{'yes':1,'no':2}}
df=df.replace(A)
df
```

Out[134]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	1	southwest	16884.92400
1	18	male	33.770	1	2	southeast	1725.55230
2	28	male	33.000	3	2	southeast	4449.46200
3	33	male	22.705	0	2	northwest	21984.47061
4	32	male	28.880	0	2	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	male	30.970	3	2	northwest	10600.54830
1334	18	female	31.920	0	2	northeast	2205.98080
1335	18	female	36.850	0	2	southeast	1629.83350
1336	21	female	25.800	0	2	southwest	2007.94500
1337	61	female	29.070	0	1	northwest	29141.36030

1338 rows × 7 columns

```
In [135]: s={"sex":{"male":1,'female':2}}
df=df.replace(s)
df
```

Out[135]:

	age	sex	bmi	children	smoker	region	charges
0	19	2	27.900	0	1	southwest	16884.92400
1	18	1	33.770	1	2	southeast	1725.55230
2	28	1	33.000	3	2	southeast	4449.46200
3	33	1	22.705	0	2	northwest	21984.47061
4	32	1	28.880	0	2	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	1	30.970	3	2	northwest	10600.54830
1334	18	2	31.920	0	2	northeast	2205.98080
1335	18	2	36.850	0	2	southeast	1629.83350
1336	21	2	25.800	0	2	southwest	2007.94500
1337	61	2	29.070	0	1	northwest	29141.36030

1338 rows × 7 columns

```
In [136]: r={"region":{"southwest":1,'southeast':2,'northwest':3,'northeast':4}}
df=df.replace(r)
df
```

Out[136]:

	age	sex	bmi	children	smoker	region	charges
0	19	2	27.900	0	1	1	16884.92400
1	18	1	33.770	1	2	2	1725.55230
2	28	1	33.000	3	2	2	4449.46200
3	33	1	22.705	0	2	3	21984.47061
4	32	1	28.880	0	2	3	3866.85520
...	...	...	...	...	...	...	...
1333	50	1	30.970	3	2	3	10600.54830
1334	18	2	31.920	0	2	4	2205.98080
1335	18	2	36.850	0	2	2	1629.83350
1336	21	2	25.800	0	2	1	2007.94500
1337	61	2	29.070	0	1	3	29141.36030

1338 rows × 7 columns

```
In [137]: #2.TO FIND DIMENSIONS OF X and y
features = df.columns[0:2]
target = df.columns[-1]
X = df[features].values
y = df[target].values
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=17)
print("The dimension of X_train is {}".format(X_train.shape))
print("The dimension of X_test is {}".format(X_test.shape))
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

The dimension of X\_train is (936, 2)  
The dimension of X\_test is (402, 2)

```
In [140]: #3.TO FIND SCORE OF LinearRegression()
lir = LinearRegression()
lir.fit(X_train, y_train)
lir.score(X_test,y_test)
print(lir.score(X_test,y_test))
#4.TO FIND TRAIN AND TEST SCORE OF LinearRegression()
train_score_lir = lir.score(X_train, y_train)
test_score_lir = lir.score(X_test, y_test)
print("\nLinear Regression Model:\n")
print("The train score for lir model is {}".format(train_score_lir))
print("The test score for lir model is {}".format(test_score_lir))
```

0.1022033122179885


Linear Regression Model:

The train score for lir model is 0.08144731818197626

The test score for lir model is 0.1022033122179885

graph of linear regression:

```
In [15]: plt.scatter(x_test,y_test,color='b')
plt.scatter(x_train,y_train,color='k')
plt.show()
```




**In the case of smaller dataset Taken(\*not necessary)**

```
In [16]: df500=df[:][:500]
```

```
In [17]: df500.fillna(method='ffill',inplace=True)
x=np.array(df500['age']).reshape(-1,1)
y=np.array(df500['charges']).reshape(-1,1)
df500.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
regr=LinearRegression()
regr.fit(x_train,y_train)
print("Regression:",regr.score(x_test,y_test))
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='r')
plt.plot(x_test,y_pred,color='k')
plt.show()
```

Regression: 0.1293688174041372



## RIDGE REGRESSION

```
In [141]: #2.TO FIND DIMENSIONS OF X and y
features = df.columns[0:2]
target = df.columns[-1]
#X and y values
x = df[features].values
y = df[target].values
#split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=17)
print("The dimension of X_train is {}".format(x_train.shape))
print("The dimension of X_test is {}".format(x_test.shape))
#Scale features
scaler = StandardScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.transform(x_test)
```


The dimension of X\_train is (936, 2)  
The dimension of X\_test is (402, 2)

```
In [33]: ridgeReg = Ridge(alpha=10)
ridgeReg.fit(X_train,y_train)
Print(ridgeReg.score(X_test, y_test))
train_score_ridge = ridgeReg.score(X_train, y_train)
test_score_ridge = ridgeReg.score(X_test, y_test)
print("\nRidge Model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.08143796463046804  
The test score for ridge model is 0.10202509697425632

```
In [34]: plt.figure(figsize = (10, 10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge;  $\alpha = 10$ ')
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression')
plt.xticks(rotation = 90)
plt.legend()
plt.show()
```



## LASSO REGRESSION

```
In [36]: lr=Lasso(alpha=10)
lr.fit(x_train,y_train)
print(lr.score(x_test,y_test))
```

0.10226046691368151



```
In [37]: print("\nLasso Model: \n")
lr = Lasso(alpha = 10)
lr.fit(X_train,y_train)
train_score_ls =lr.score(X_train,y_train)
test_score_ls =lr.score(X_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))
```

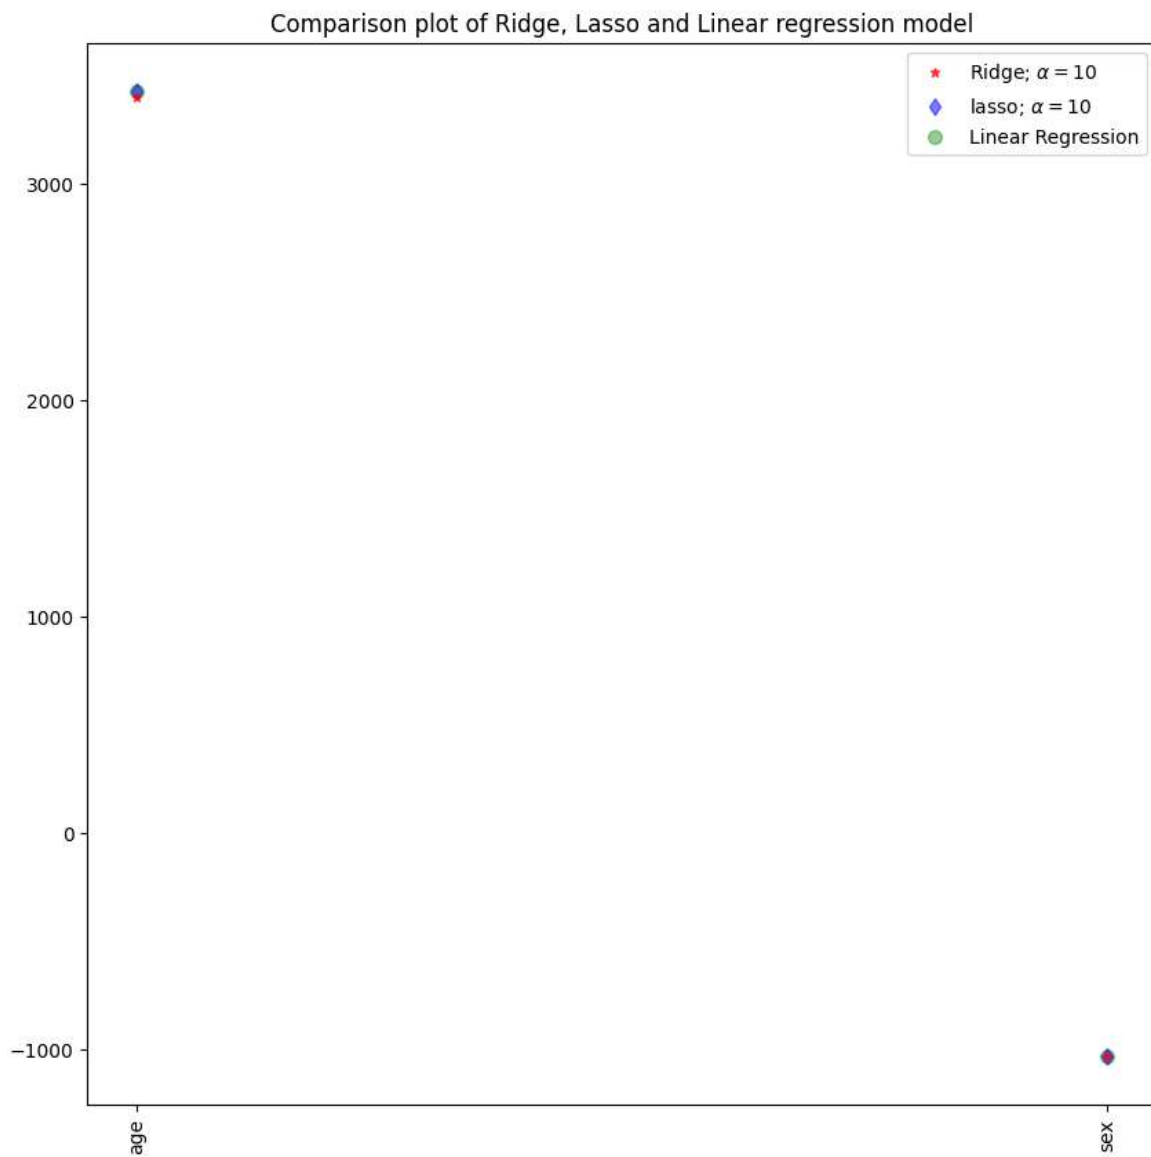
Lasso Model:

The train score for ls model is 0.08144600503720834  
The test score for ls model is 0.10226046691368151

```
In [44]: from sklearn.linear_model import LassoCV
#Lasso Cross validation
lasso_cv = LassoCV(alphas = [0.0001, 0.001,0.01, 0.1, 1, 10], random_state=0).fit(X_train, y_train)
#score
print(lasso_cv.score(X_train, y_train))
print(lasso_cv.score(X_test, y_test))
```


0.08144600503720834  
0.10226046691368151

```
In [45]: plt.figure(figsize = (10, 10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge;  $\alpha=10$ ',zorder=6)
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'lasso;  $\alpha=10$ ',zorder=6)
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression')
plt.xticks(rotation = 90)
plt.legend()
plt.title("Comparison plot of Ridge, Lasso and Linear regression model")
plt.show()
```



```
In [47]: pd.Series(lr.coef_, features).sort_values(ascending = True).plot(kind = "bar")
```

Out[47]: <Axes: >



## ELASTIC NET

```
In [48]: from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(x,y)
print(regr.coef_)
print(regr.intercept_)
regr.score(x,y)
```

```
[ 257.44760657 -511.96287073]
3941.933287820537
```

Out[48]: 0.09164498902013407

Here ridge and lasso regression have low accuracy. So, Elastic Net also has low accuracy which depends on the values of ridge and lasso.

```
In [49]: y_pred_elastic=regr.predict(x_train)
```

```
In [50]: mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
print("Mean Squared Error on test set",mean_squared_error)
```

Mean Squared Error on test set 251337238.06352752

## 5.Data Prediction&Evaluation

```
In [51]: prediction=lr.predict(x_test)
```

```
In [52]: from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2_score:",r2)
```

R2\_score: 0.1022033122179885

## To find Error

```
In [53]: print('MAE:',metrics.mean_absolute_error(y_test,prediction))
print('MSE:',metrics.mean_squared_error(y_test,prediction))
print('RMSE:',np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

MAE: 8564.838037174472  
MSE: 109441886.11907542  
RMSE: 10461.447611065852

## Cross Validation

```
In [54]: ridge_cv=RidgeCV(alphas=[1,10,100]).fit(X_train,y_train)
print("The train score for ridgecrossvalidation model is : {}".format(ridge_cv.score(X_train,y_train)))
print("The test score for ridgecrossvalidation model is :{}".format(ridge_cv.score(X_test,y_test)))
```

The train score for ridgecrossvalidation model is : 0.08143796463046815  
The test score for ridgecrossvalidation model is :0.10202509697425854

```
In [55]: lasso_cv=LassoCV(alphas=[1,10,100]).fit(X_train,y_train)
print(" The train score for lassocrossvalidation model is : {}".format(lasso_cv.score(X_train,y_train)))
print("The test score for lassocrossvalidation model is:{}".format(lasso_cv.score(X_test,y_test)))
```

The train score for lassocrossvalidation model is : 0.08144600503720834  
The test score for lassocrossvalidation model is:0.10226046691368151

**Here the regression score(accuracy) is low.So,check the categorical models**

## LOGISTIC REGRESSION

```
In [59]: A={"smoker":{"yes":1,'no':2}}
df=df.replace(A)
df
```

Out[59]:

	age	sex	bmi	children	smoker	region	charges
0	19	2	27.900	0	1	southwest	16884.92400
1	18	1	33.770	1	2	southeast	1725.55230
2	28	1	33.000	3	2	southeast	4449.46200
3	33	1	22.705	0	2	northwest	21984.47061
4	32	1	28.880	0	2	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	1	30.970	3	2	northwest	10600.54830
1334	18	2	31.920	0	2	northeast	2205.98080
1335	18	2	36.850	0	2	southeast	1629.83350
1336	21	2	25.800	0	2	southwest	2007.94500
1337	61	2	29.070	0	1	northwest	29141.36030

1338 rows × 7 columns

```
In [60]: x=df.iloc[:, :-1].values
y=df.iloc[:, 1].values
```

```
In [62]: features_matrix=df.iloc[:, 0:4]
target_vector=df.iloc[:, -3]
print('The Features Matrix Has %d Rows And %d Columns(s)'%(features_matrix.shape))
print('The Features Matrix Has %d Rows And %d Columns(s)'%(np.array(target_vector).reshape(-1,1).shape))
```

The Features Matrix Has 1338 Rows And 4 Columns(s)  
The Features Matrix Has 1338 Rows And 1 Columns(s)

```
In [63]: features_matrix_standardized=StandardScaler().fit_transform(features_matrix)
```

```
In [64]: algorithm=LogisticRegression(max_iter=100000)
```

```
In [65]: Logistic_Regression_Model=algorithm.fit(features_matrix_standardized,target_vector)
```

```
In [66]: observation=[[0,0.01,0.0003,0.0004]]
```

```
In [67]: predictions=Logistic_Regression_Model.predict(observation)
```

```
In [68]: print('The Model Predicted The Obsevation To Belong To Class %s'%(predictions))
```

The Model Predicted The Obsevation To Belong To Class [2]

```
In [75]: features = df.columns[0:2]
target = df.columns[-2]
#X and y values
X = df[features].values
y = df[target].values
#split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=17)
print("The dimension of X_train is {}".format(X_train.shape))
print("The dimension of X_test is {}".format(X_test.shape))
#Scale features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

The dimension of X\_train is (936, 2)  
The dimension of X\_test is (402, 2)

```
In [78]: #Model
legr = LogisticRegression()
#Fit model
legr.fit(X_train, y_train)
#predict
#prediction = lr.predict(X_test)
#actual
actual = y_test
train_score_legr = legr.score(X_train, y_train)
test_score_legr = legr.score(X_test, y_test)
print("\nlogistic Model:\n")
print("The train score for legr model is {}".format(train_score_legr))
print("The test score for legr model is {}".format(test_score_legr))
```

logistic Model:

The train score for legr model is 0.2724358974358974  
The test score for legr model is 0.24129353233830847

```
In [79]: x=np.array(df['age']).reshape(-1,1)
y=np.array(df['smoker']).reshape(-1,1)
```

```
In [80]: X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.30)
lerg=LogisticRegression()
lerg.fit(X_train,y_train)
print(lerg.score(X_test,y_test))
```

0.8233830845771144

C:\Users\HP\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\utils\validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using.ravel().

y = column\_or\_1d(y, warn=True)

we can apply another regression models also...

## Decision Tree Classifier

```
In [81]: from sklearn.tree import DecisionTreeClassifier
```

```
In [82]: df['sex'].value_counts()
```

```
Out[82]: sex
1      676
2      662
Name: count, dtype: int64
```

```
In [89]: df['bmi'].value_counts()
```

```
Out[89]: bmi
32.300    13
28.310     9
30.495     8
30.875     8
31.350     8
..
46.200     1
23.800     1
44.770     1
32.120     1
30.970     1
Name: count, Length: 548, dtype: int64
```

```
In [90]: X=["sex","age","bmi"]
y=["1","2"]
all_inputs=df[X]
all_classes=df["region"]
```

```
In [91]: (X_train,X_test,y_train,y_test)=train_test_split(all_inputs,all_classes,test_size=0.5)
```

```
In [92]: clf=DecisionTreeClassifier(random_state=0)
```

```
In [93]: clf.fit(X_train,y_train)
```

```
Out[93]: DecisionTreeClassifier
DecisionTreeClassifier(random_state=0)
```

```
In [94]: score=clf.score(X_test,y_test)
print(score)
```

```
0.3034379671150972
```

## CONCULSION:

```
In [ ]: It is concluded that regression models:linear,Ridge,Lasso have 10.22% of accuracy,DecisionTreeClassifier has 30% accuracy,
logistic regression has 82% accuracy.So,Logistic regression comparatively high with another models.So,Logistic regression is the
best fit for health insurance.
```